

Generative Engine Platform

V2 release – general communication

We've been cooking

As promised, we will upgrade tonight our beloved **Generative Engine Platform** to its **V2**, less than 1 year after initial launch!

Since then we've seen tremendous traction from across Capgemini, with thousands of active users, billions of tokens exchanged, dozens of assets built, and a great number of success stories from the engineering communities supporting our clients!

This upcoming release is massive!

We've been focusing on scale-up, collaboration, and of course Agentic AI!

-  **Refreshed UI overall**, introduction of collaborative studios, including unlimited plan, an important milestone where Generative Engine Platform will remain free for all, and some power users or teams will get access to more capabilities through a paid tier.
-  **OpenAI compatible APIs** means engineers can build agentic applications using their preferred frameworks and tools on their side.
-  **New models!** We are adding Reasoning Models with Gemini 2.0, DeepSeek-R1 and o3-mini models, as well as Claude Sonnet 3.7, best model in the world for software engineers. We also add Perplexity Sonar for agentic « deep search » across the web.
-  **Assistants promoted to AI Agents!** We continuously upgrade them with better models (Lyn is now powered by Sonnet 3.7!), refined instructions, and new agentic capabilities to achieve more complex tasks. You can now search in ArXiv database, run data analytics and visualization, and more.
-  **Introduction of *Marketplace*** our new page empowering anyone to share their AI work with everyone else! We start with 10 awesome contributions and we look forward to getting a lot more from engineers soon!

The release is planned for **Tuesday 25th at around 8pm CET** – still subject to last minute change.

We'll do a more comprehensive communication to users and communities in coming weeks, while we start working on v2.1.

Happy generating 

V2 release – detailed change log

Overall UX/UI and concept changes:

- New look&feel for the whole platform with simplified interface
- Concept rework: You can now create multiple studios. A studio is a segregated space and can be individual (personal) or shared (collaborative). You can host multiple Knowledge Bases (previous workspaces) into a studio. API keys will be bound to each studio (vs users).
- New pop-up mechanism to display an information or video about release changes
- New introduction and changelog videos, featuring Jenn, your AI business partner
- 3 new videos to explain how to use the studio
- New changelog display and dedicated page to give you more information
- Chat History has been greatly improved:
 - Chat history is now segregated to ensure better data control – you now have a chat history for agents, and another chat history for each studio
 - When creating a new conversation, the title of the conversation (in the history) will be summarized based on your first 2 messages. The title can also be modified manually.
 - You can now search from previous conversation (search on title and content)
 - You can now see the date of last modification of a conversation
 - By default, conversations will now be sorted by last modification, to make active conversation remain on the first page

New models:

- 3 LLMs
 - Claude Sonnet 3.7
 - Gemini 2.0 Flash Lite
 - Perplexity Sonar
- 3 New Reasoning models:
 - Gemini-2.0 Flash Thinking
 - DeepSeek R1
 - o3-mini

Marketplace:

- Collection is now a marketplace where will be hosted many AI initiatives from the community. Feel free to reach out if you want to share an asset you have built !
- New tag, filtering and search mechanisms have been implemented in the marketplace

AI Agents:

- Assistants have been promoted to Agents 😊 This new denomination should better reflect their growing capacity and our vision to enhance them in the near future
- The AI Team catalog is now available to browse through Agents
- You can now favorite your preferred agents for quick access on the left menu
- Each agent now proposes Tasks examples to guide you on what each Agent can do, and help create more efficient prompts
- You can now create a new conversation with the + sign, available at all time
- We are introducing agentic capabilities, available by typing "/" at the beginning of your prompt. We are releasing 3 new and experimental capabilities:
 - A capability to analyze your Excel or Data file and provide you with an analysis or metrics => can we highlight this as a task on Daniel - the Data & AI Engineer

- A capability that can search and analyze arXiv research papers => can we highlight this one on Roger - the Research Engineer ?
- A capability that can create diagrams based on description or image => can we highlight this one on Daniel as well ?
- A new AI Agent "Roger" is now available to search information on the Internet and create reports for you leveraging Perplexity agentic search model (Sonar)
- Several Agents have been upgraded to newer generation models (ex: Lyn - the Software Engineer - to Sonnet 3.7 the best coding model in the world)

Studio:

- Rework to make the studio a cornerstone entity. It now holds the API Key management, User management, Knowledge Base management, and chat history management.
- You can now create multiple studios amongst 3 tiers:
 - Personal Free studio : you can only have 1 and it can be used for experimenting. Most features and models are available, but with a limited weekly usage
 - Personal Paid studio : you can have multiple ones and they can be used for building or heavy personal usage. They will be billed a small fixed monthly fee. Weekly usage is greatly increased compared to the free studio.
 - Collaborative studio: you can have multiple ones and they can be used for working as a team on a project or building a production-grade asset. Multiple users can use this pay-as-you-go studio and you can generate a project API key for your asset. All models and features are available in this studio, and usage is close to unlimited
- You can now select a profile (standard, dev, data scientist or full) to adapt the features displayed on the interface.
- You can now create ephemeral conversations in the Playground. Ephemeral conversation will not be stored on our platform and will not appear on your history.
- You can now only create private Knowledge Bases and allow specific users or specific studios to query it.
- A collaborative studio can invite users with specific rights: owner, contributor, or consumer. These rights apply to all knowledge bases of this studio.
- Knowledge bases are now limiting the amount of vectorized data to ensure cost controls.
- Limits can be fine-tuned per studio, to accommodate to specific use cases.
- A new endpoint has been created to provide OpenAI compatibility. Please check the updated documentation for more details on how to use it. It enables agentic and tools use capability to the API when building your assets.

Technical:

- Huge adapter rework, allowing us to support wider features on more models
- Technical improvement regarding the management of our legacy users
- Updated langchain to the latest version
- New instantiation for a more autonomous authentication/authorization mechanism
- Creation of a new instantiation of the platform for better scalability and flexibility
- Migration of selected data
- Improved our observability stack
- Improved technical and cyber documentation
- Improved life cycle management
- Industrialization of our delivery pipeline
- Various bug fixes